

Towards Modular Compilation using Higher-order Effects

Anonymous author

Anonymous affiliation

Abstract

Compilers transform a human readable source language into machine readable target language. Nanopass compilers simplify this approach by breaking up this transformation into small steps that are more understandable, maintainable, and extensible. We propose a semantics-driven variant of the nanopass compiler architecture exploring the use of effects and handlers to model the intermediate languages and the transformation passes, respectively. Our approach is fully typed and ensures that all cases in the compiler are covered. Additionally, by using an effect system we abstract over the control flow of the intermediate language making the compiler even more flexible. We apply this approach to a minimal compiler from a language with arithmetic and let-bound variables to a minimal subset of X86 machine code with variables. In the future, we hope to extend this work to compile a larger and more complicated language and we envision a formal verification framework from compilers written in this style.

2012 ACM Subject Classification Theory of computation → Program semantics; Software and its engineering → Compilers

Keywords and phrases algebraic effects and handlers, higher-order effects, monadic semantics, modularity, compilation, nanopass

Digital Object Identifier 10.4230/OASIS.CVIT.2016.23

1 Introduction

The essence of a compiler is a function from a source language, suitable for humans, to a machine language, suitable for computers. As our computers have become more powerful we have seen increasingly complex compilers providing extensive safety guarantees and powerful optimizations. To manage this complexity, modern compilers are designed as a composition of multiple passes. However, the total number of passes has traditionally been kept low for the sake of performance, because each pass adds extra overhead. Thus, compiler passes are more complicated than necessary and therefore harder to *understand*, *maintain*, and *extend*.

To address this problem, Sarkar, Waddell, and Dybvig [7] introduce the nanopass compiler architecture. In the nanopass architecture, each pass is designed to be as simple as possible. It is not a problem to use many more passes than are used in traditional compilers. To address concerns about the performance of this architecture, Keep and Dybvig [4] show that it is possible to write a competitive commercial compiler using this architecture.

While the development of the nanopass architecture and the development of a commercial compiler is a great engineering achievement, we believe the theoretical foundation is underexplored. In this paper, we present our ongoing work on developing a semantics-driven nanopass compiler architecture. Our approach has the practical advantage of preventing type errors in the compiler and ensuring all cases are covered. Additionally, our approach abstracts over the control flow giving many of the benefits of continuation-passing-style while retaining a simple monadic interface. During the development of our approach, we are anticipating the possibility of verifying the correctness of compilers written in this style.

Concretely, we make the following contributions:

- We introduce a novel approach to designing and implementing practical compilers while staying close to formal denotational semantics specifications (Section 2).

■ We demonstrate our approach on a simple language with arithmetic and let-bound variables (Section 2). We compile this language to a subset of X86 with variables. We are still working a method to allocate these variables on the stack.

2 Compiling with Higher-order Effects

In this section, we present our approach by applying it to a very simple language with arithmetic and let-bound variables. The target language of our compiler is X86 machine code. We explain the required concepts as we develop our compiler for this language. The specifications and compiler passes are presented in a simplified notation, but we have implemented all the work we present here in the Agda programming language [2].

We start off by assuming our parser and possibly type checker has finished and produced an abstract syntax tree which follows the grammar described in Figure 1. Our language has integers, addition, subtraction, negation, a read operation to read an input integer, and a let to bind variables and a var to refer to bound variables. The abstract syntax is unusual because we reuse the variable binding facilities from our host language in the form of the λx . binding in the let constructor. This style of abstract syntax is called parametric higher-order abstract syntax (PHOAS) [3]. It allows us to avoid the complexities of variable binding and thus simplify our presentation.

$$\begin{aligned} \text{expr} ::= & \text{int}(n) \\ & | \text{add}(\text{expr}, \text{expr}) \\ & | \text{sub}(\text{expr}, \text{expr}) \\ & | \text{neg}(\text{expr}) \\ & | \text{read} \\ & | \text{let}(\text{expr}, \lambda x. \text{expr}) \\ & | \text{var}(x) \end{aligned}$$

■ **Figure 1** Abstract syntax of our simple language with arithmetic and let-bound variables.

The first step of our compilation pipeline will be to denote these syntactic constructs onto a set of semantic algebraic operations in the sense of algebraic effects [5]. As is customary when using algebraic effects in functional programming languages, we group these operations under units we call effects. We could group every operation under a single effect, however with the benefit of hindsight we decide to distribute the operations over four effects: Int, Arith, Read, and Let.

Figure 2 shows the operations that correspond directly to our source language. However, there are some particularities we address individually:

- To keep our example simple, we have chosen to use this single type for all values, but we keep the type abstract and simply call it ‘val’. It is crucial to keep this type abstract for reasons we explain at the end of this section. Also note that we now need a special *int* operation to inject integers into this abstract value type.
- We write the surrounding monadic context as m , which provides the standard $\gg=$, \gg , and *return* operations. The monadic context of each operation always includes at least the effect that the operation belongs to, but it can accomodate other effects too. For

$$\begin{array}{ll}
\text{int} : \mathbb{Z} \rightarrow m \text{ val} & \\
& \text{read} : m \text{ val} \\
\\
\text{add} : \text{val} \rightarrow \text{val} \rightarrow m \text{ val} & \\
\text{sub} : \text{val} \rightarrow \text{val} \rightarrow m \text{ val} & \text{let} : m \text{ val} \rightarrow (\text{val} \rightarrow m \text{ val}) \rightarrow m \text{ val} \\
\text{neg} : \text{val} \rightarrow m \text{ val} &
\end{array}$$

■ **Figure 2** The operations of our source language with their signatures. Divided into the following effects from left to right and then top to bottom: Int, Read, Arith, and Let.

78 example, the monadic context of $\text{int } 1 \gg \lambda x. \text{add } x \ x$ contains at least the Int and Arith
 79 effects.

80 ■ Readers knowledgeable about effect systems might notice that the *let* operation has
 81 arguments that are themselves monadic computations. In standard algebraic effects
 82 and handlers this is not allowed, however our approach uses a novel higher-order effect
 83 formalism that does support such effectful subcomputations [1].

84 The first pass of our compiler pipeline is to map our abstract syntax from Figure 1 onto
 85 the operations we have defined for our source language from Figure 2. This mapping, called
 86 a denotation and written using the $\llbracket \cdot \rrbracket$ notation, is a recursive traversal of the abstract syntax
 87 tree shown in Figure 3. The result of this mapping is a monadic computation involving the
 88 Int, Arith, Read, and Let effects.

$$\begin{aligned}
\llbracket \text{int}(n) \rrbracket &= \text{int } n \\
\llbracket \text{add}(e_1, e_2) \rrbracket &= \llbracket e_1 \rrbracket \gg \lambda x. \llbracket e_2 \rrbracket \gg \lambda y. \text{add } x \ y \\
\llbracket \text{sub}(e_1, e_2) \rrbracket &= \llbracket e_1 \rrbracket \gg \lambda x. \llbracket e_2 \rrbracket \gg \lambda y. \text{sub } x \ y \\
\llbracket \text{neg}(e) \rrbracket &= \llbracket e \rrbracket \gg \lambda x. \text{neg } x \\
\llbracket \text{read} \rrbracket &= \text{read} \\
\llbracket \text{let}(e, f) \rrbracket &= \text{let } \llbracket e \rrbracket (\lambda x. \llbracket f \rrbracket x) \\
\llbracket \text{var}(x) \rrbracket &= \text{return } x
\end{aligned}$$

■ **Figure 3** Denotational mapping from our abstract syntax onto our initial set of effectful operations.

89 Now that we have denoted our syntactic elements into our semantic domain as operations,
 90 we can start refining these operations to get closer to the desired target language which is
 91 X86 in our case. In the practice of algebraic effects, this refinement is facilitated by handlers
 92 as introduced by Plotkin and Pretnar [6]. These handlers give us access to the operations
 93 that occur in the program and the continuation of the program which we will simply call k .
 94 We only have to provide the rules that map the operation and continuation onto our semantic
 95 domain consisting of primitives, existing operations, or newly introduced operations.

96 The effect that we choose to handle first is the Let effect, which only has the *let* operation.
 97 We handle this operation by running the right hand side of the binding, passing the resulting
 98 value to the body, and finally passing the result of that to the continuation. In code that

23:4 Towards Modular Compilation using Higher-order Effects

99 looks as follows:

100 handle (let $e\ f$) $k = e \gg \lambda x. f\ x \gg \lambda z. k\ z$
 101

102 Note that this defines a strict semantics for our let bindings. By using a different handler
 103 we could give different semantics to our language. This is an example of the flexibility of
 104 algebraic effects and handlers.

$addq : val \rightarrow val \rightarrow m\ ()$	
$subq : val \rightarrow val \rightarrow m\ ()$	
$negq : val \rightarrow m\ ()$	
$movq : val \rightarrow val \rightarrow m\ ()$	$x86var : m\ val$
$callq : lab \rightarrow m\ ()$	
$reg : Register \rightarrow m\ val$	

■ **Figure 4** X86 related effects. On the left X86. On the right X86Var.

105 At this point, since our language is so simple we can already begin translating into
 106 our target language. In Figure 4 we show the a minimal subset of X86 that we need to
 107 compile our Arith and Read effects. This subset contains in-place arithmetic instructions,
 108 the ubiquitous move instruction, the call instruction, and an operation to inject concrete
 109 registers into our abstract value type. Additionally, we add an operation to generate fresh
 110 variables and inject them into our abstract value type.

111 We can translate our Arith effect operations into X86 operations by creating a fresh X86
 112 variable, populating it, and then applying the in-place arithmetic operation to the variable.
 113 So, we write handler as follows:

114 handle (add $x\ y$) $k = x86var \gg \lambda z. movq\ x\ z \gg addq\ y\ z \gg k\ z$
 115 handle (sub $x\ y$) $k = x86var \gg \lambda z. movq\ x\ z \gg subq\ y\ z \gg k\ z$
 116 handle (neg x) $k = x86var \gg \lambda z. movq\ x\ z \gg negq\ z \gg k\ z$
 117

118 The *read* operation requires us to call a function that we will assume is defined in a standard
 119 library called `read_int`. This function places its output in the `%rax` register, so we have to
 120 move it to avoid it being overwritten by other parts of our program. The full definition of
 121 our handler for the Read effect is as follows:

122 handle read $k = x86var \gg \lambda z. callq\ read_int \gg reg\ \%rax \gg \lambda x. movq\ x\ z \gg k\ z$
 123

124 The final challenge to complete this minimal compiler is to allocate the X86 variables
 125 on the stack. Conceptually, this requires us to give each *x86var* operation and give each its
 126 own location of the stack. Keeping track of such information in our handler, however, is
 127 something we have not yet needed to do for the passes up to this point. The main problem is
 128 that we need to pass a valid value as input to the continuation of each operation to get access
 129 to the operations in the continuation. When the type of the argument of the continuation is a
 130 unit type, such as for the X86 arithmetic, move, and call instructions, then we can simply
 131 construct the unit value and pass that to the continuation. However, if we allow arbitrary
 132 concrete types as results of our operations then there is no guarantee that we can construct a
 133 value of that type and it is not always possible to reconstruct the higher-order representation
 134 if we do manage to construct such a value.

135 Luckily, we have chosen to keep our value type abstract. So, we still have a choice to
136 instantiate it to a type that suits our purpose. Instantiating it to the unit type would make
137 it possible for us to construct a value to pass to the continuation. However, we would no
138 longer be able to distinguish values passed to different continuations, so it would not be
139 possible to reconstruct a higher order representation for further manipulation. Instead, our
140 solution is to instantiate it to the type of natural numbers and to pass a unique number to
141 each continuation. Whenever we encounter such a natural number in the rest of the program,
142 we know which continuation it originated from. Hence, we are able to reconstruct a higher
143 order operation.

144 This process sounds complicated, but we expect it is possible to expose the ability to
145 keep track of the required information through an easy to use API.

146 **3 Related Work**

147 The origins of this work can be traced back to Eelco Visser's work on the Spoofox Language
148 Workbench [9]. As part of Spoofox, Eelco is one of the designers of the Stratego [8] program
149 transformation language. While Stratego can be used for developing compilers, Eelco was
150 still looking for a way of specifying compilers that also abstracts over the control flow. We
151 hope that this work can be the start of an answer to that research direction.

152 As mentioned in the introduction, our approach embraces the nanopass architecture [7, 4].
153 We improve upon this work by putting it on more formal foundations, making it fully typed
154 to prevent common errors and even check that all cases are covered, and abstracting over
155 the control flow in the compiler.

156 Our semantics-driven approach using an effect system is inspired by the work on symbolic
157 execution by Wei et al. [10].

158 For our vision on verification of compilers we have taken inspiration from Interaction
159 Trees [11]. Interaction Trees use algebraic laws of effects to prove the correctness of compilers.
160 We hope to learn from that technique to prove the correctness of compilers written using our
161 approach.

162 **4 Conclusions and Future Work**

163 We have presented a new semantics-driven approach to writing compilers by using effect
164 operations as an intermediate representation. We use effect handlers to iteratively refine
165 operations in terms of increasingly lower level operations to finally reach a target machine
166 language.

167 We have shown a concrete example of this approach applied to a very simple language
168 with arithmetic and let-bound variables. We show the implementation of a denotation
169 function and handlers which compile this language is compiled in several passes to X86
170 machine language with variables. Currently, we are working on a stack allocation pass to
171 complete this minimal compiler.

172 In the future, we would like to complete the minimal compiler and extend it with more
173 complicated language constructs such as conditionals and anonymous functions. Furthermore,
174 we would like to explore the verification of our compilers using algebraic laws for our effect
175 operations.

References

- 177 **1** Casper Bach Poulsen and Cas van der Rest. Hefty algebras – modular elaboration for
178 higher-order algebraic operations. Under submission, 2022.
- 179 **2** Ana Bove, Peter Dybjer, and Ulf Norell. A brief overview of agda – a functional language
180 with dependent types. In Stefan Berghofer, Tobias Nipkow, Christian Urban, and Makarius
181 Wenzel, editors, *Theorem Proving in Higher Order Logics*, pages 73–78, Berlin, Heidelberg,
182 2009. Springer Berlin Heidelberg.
- 183 **3** Adam Chlipala. Parametric higher-order abstract syntax for mechanized semantics. *SIGPLAN*
184 *Not.*, 43(9):143–156, sep 2008. URL: [https://doi-org.tudelft.idm.oclc.org/10.1145/](https://doi-org.tudelft.idm.oclc.org/10.1145/1411203.1411226)
185 [1411203.1411226](https://doi-org.tudelft.idm.oclc.org/10.1145/1411203.1411226), doi:10.1145/1411203.1411226.
- 186 **4** Andrew W. Keep and R. Kent Dybvig. A nanopass framework for commercial compiler
187 development. *SIGPLAN Not.*, 48(9):343–350, sep 2013. URL: [https://doi-org.tudelft.](https://doi-org.tudelft.idm.oclc.org/10.1145/2544174.2500618)
188 [idm.oclc.org/10.1145/2544174.2500618](https://doi-org.tudelft.idm.oclc.org/10.1145/2544174.2500618), doi:10.1145/2544174.2500618.
- 189 **5** Gordon Plotkin and John Power. Adequacy for algebraic effects. In Furio Honsell and Marino
190 Miculan, editors, *Foundations of Software Science and Computation Structures*, pages 1–24,
191 Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- 192 **6** Gordon Plotkin and Matija Pretnar. Handlers of algebraic effects. In Giuseppe Castagna,
193 editor, *Programming Languages and Systems*, pages 80–94, Berlin, Heidelberg, 2009. Springer
194 Berlin Heidelberg.
- 195 **7** Dipanwita Sarkar, Oscar Waddell, and R. Kent Dybvig. A nanopass infrastructure for
196 compiler education. In *Proceedings of the Ninth ACM SIGPLAN International Conference on*
197 *Functional Programming*, ICFP ’04, page 201–212, New York, NY, USA, 2004. Association for
198 Computing Machinery. URL: [https://doi-org.tudelft.idm.oclc.org/10.1145/1016850.](https://doi-org.tudelft.idm.oclc.org/10.1145/1016850.1016878)
199 [1016878](https://doi-org.tudelft.idm.oclc.org/10.1145/1016850.1016878), doi:10.1145/1016850.1016878.
- 200 **8** Eelco Visser. Stratego: A language for program transformation based on rewriting strategies
201 system description of stratego 0.5. In Aart Middeldorp, editor, *Rewriting Techniques and*
202 *Applications*, pages 357–361, Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- 203 **9** Guido H. Wachsmuth, Gabriël D.P. Konat, and Eelco Visser. Language design with the
204 spoofax language workbench. *IEEE Software*, 31(5):35–43, 2014. doi:10.1109/MS.2014.100.
- 205 **10** Guannan Wei, Oliver Bračevac, Shangyin Tan, and Tiark Rompf. Compiling symbolic execution
206 with staging and algebraic effects. *Proc. ACM Program. Lang.*, 4(OOPSLA), nov 2020. URL:
207 <https://doi-org.tudelft.idm.oclc.org/10.1145/3428232>, doi:10.1145/3428232.
- 208 **11** Li-yao Xia, Yannick Zakowski, Paul He, Chung-Kil Hur, Gregory Malecha, Benjamin C. Pierce,
209 and Steve Zdancewic. Interaction trees: Representing recursive and impure programs in coq.
210 *Proc. ACM Program. Lang.*, 4(POPL), dec 2019. URL: [https://doi-org.tudelft.idm.oclc.](https://doi-org.tudelft.idm.oclc.org/10.1145/3371119)
211 [org/10.1145/3371119](https://doi-org.tudelft.idm.oclc.org/10.1145/3371119), doi:10.1145/3371119.