

# Towards Modular Compilation using Higher-order Effects

Anonymous author

Anonymous affiliation

## Abstract

Compilers transform a human readable source language into machine readable target language. Nanopass compilers simplify this approach by breaking up this transformation into small steps that are more understandable, maintainable, and extensible. We propose a semantics-driven variant of the nanopass compiler architecture exploring the use of effects and handlers to model the intermediate languages and the transformation passes, respectively. Our approach is fully typed and ensures that all cases in the compiler are covered. Additionally, by using an effect system we abstract over the control flow of the intermediate language making the compiler even more flexible. We apply this approach to a minimal compiler from a language with arithmetic and let-bound variables to a string of pretty printed X86 instructions. In the future, we hope to extend this work to compile a larger and more complicated language and we envision a formal verification framework from compilers written in this style.

**2012 ACM Subject Classification** Theory of computation → Program semantics; Software and its engineering → Compilers

**Keywords and phrases** algebraic effects and handlers, higher-order effects, monadic semantics, modularity, compilation, nanopass

**Digital Object Identifier** 10.4230/OASICS.CVIT.2016.23

## 1 Introduction

The essence of a compiler is a function from a source language, suitable for humans, to a machine language, suitable for computers. As our computers have become more powerful we have seen increasingly complex compilers providing extensive safety guarantees and powerful optimizations. To manage this complexity, modern compilers are designed as a composition of multiple passes. However, the total number of passes has traditionally been kept low for the sake of performance, because each pass adds extra overhead. Thus, compiler passes are more complicated than necessary and therefore harder to *understand*, *maintain*, and *extend*.

To address this problem, Sarkar, Waddell, and Dybvig [7] introduce the nanopass compiler architecture. In the nanopass architecture, each pass is designed to be as simple as possible. It is not a problem to use many more passes than are used in traditional compilers. To address concerns about the performance of this architecture, Keep and Dybvig [4] show that it is possible to write a competitive commercial compiler using this architecture.

While the development of the nanopass architecture and the development of a commercial compiler is a great engineering achievement, we believe the theoretical foundation is underexplored. In this paper, we present our ongoing work on developing a semantics-driven nanopass compiler architecture. The foundation of our approach is a higher-order effect system [1] which is a state of the art technique for modeling the semantics of programming languages with side effects. Our approach has the practical advantage of preventing type errors in the compiler and ensuring all cases are covered. Additionally, our approach abstracts over the control flow giving many of the benefits of continuation-passing-style while retaining a simple monadic interface. While developing our approach, we are anticipating the possibility of verifying the correctness of compilers written in this style. We discuss that possibility briefly in our future work (Section 4).

- 46 Concretely, we make the following contributions:
- 47 ■ We introduce a novel approach to designing and implementing practical compilers while
  - 48 staying close to formal denotational semantics specifications (Section 2).
  - 49 ■ We demonstrate our approach on a simple language with arithmetic and let-bound
  - 50 variables by compiling it to a subset of X86 (Section 2).

## 51      2      Compiling with Higher-order Effects

52 In this section, we present our approach by applying it to a very simple language with  
 53 arithmetic and let-bound variables. The target language of our compiler is X86 machine  
 54 code. We explain the required concepts as we develop our compiler for this language.  
 55 The specifications and compiler passes are presented in a simplified notation, but we have  
 56 implemented all the work we present here in the Agda programming language [2]. Our code  
 57 can be found on GitHub<sup>1</sup>.

58 We start off by assuming our parser and possibly type checker has finished and produced  
 59 an abstract syntax tree which follows the grammar described in Figure 1. Our language has  
 60 integers, addition, subtraction, negation, a read operation to read an input integer, and a  
 61 let to bind variables and a var to refer to bound variables. The abstract syntax is unusual  
 62 because we reuse the variable binding facilities from our host language in the form of the  $\lambda x$ .  
 63 binding in the let constructor. This style of abstract syntax is called parametric higher-order  
 64 abstract syntax (PHOAS) [3]. It allows us to avoid the complexities of variable binding and  
 65 thus simplify our presentation.

$$\begin{aligned}
 \text{expr} ::= & \text{int}(n) \\
 & | \text{add}(\text{expr}, \text{expr}) \\
 & | \text{sub}(\text{expr}, \text{expr}) \\
 & | \text{neg}(\text{expr}) \\
 & | \text{read} \\
 & | \text{let}(\text{expr}, \lambda x. \text{expr}) \\
 & | \text{var}(x)
 \end{aligned}$$

■ **Figure 1** Abstract syntax of our simple language with arithmetic and let-bound variables.

66 The first step of our compilation pipeline will be to denote these syntactic constructs onto  
 67 a set of semantic algebraic operations in the sense of algebraic effects [5]. As is customary  
 68 when using algebraic effects in functional programming languages, we group these operations  
 69 under units we call effects. We could group every operation under a single effect, however  
 70 with the benefit of hindsight we decide to distribute the operations over four effects: Int,  
 71 Arith, Read, and Let.

72 Figure 2 shows the operations that correspond directly to our source language. However,  
 73 there are some particularities we address individually:

- 74 ■ To keep our example simple, we have chosen to use this single type for all values, but we
- 75 keep the type abstract and simply call it ‘val’. It is crucial to keep this type abstract for

---

<sup>1</sup> <https://github.com/heft-lang/hefty-compilation>

<b>effect Int where</b> $int : \mathbb{Z} \rightarrow m\ val$	<b>effect Read where</b> $read : m\ val$
<b>effect Arith where</b> $add : val \rightarrow val \rightarrow m\ val$ $sub : val \rightarrow val \rightarrow m\ val$ $neg : val \rightarrow m\ val$	<b>effect Let where</b> $let : m\ val \rightarrow (val \rightarrow m\ val) \rightarrow m\ val$

■ **Figure 2** The effects of our source language and their operations with type signatures.

76 reasons we explain at the end of this section. Also note that we now need a special *int*  
77 operation to inject integers into this abstract value type.

78 ■ We write the surrounding monadic context as  $m$ , which provides the standard  $\gg=$ ,  $\gg$ ,  
79 and *return* operations. The monadic context of each operation always includes at least  
80 the effect that the operation belongs to, but it can accomodate other effects too. For  
81 example, the monadic context of  $int\ 1 \gg= \lambda x. add\ x\ x$  contains at least the Int and Arith  
82 effects.

83 ■ Readers knowledgeable about effect systems might notice that the *let* operation has  
84 arguments that are themselves monadic computations. In standard algebraic effects  
85 and handlers this is not allowed, however our approach uses a novel higher-order effect  
86 formalism that does support such effectful subcomputations [1].

87 The first pass of our compiler pipeline is to map our abstract syntax from Figure 1 onto  
88 the operations we have defined for our source language from Figure 2. This mapping, called  
89 a denotation and written using the  $\llbracket \cdot \rrbracket$  notation, is a recursive traversal of the abstract syntax  
90 tree shown in Figure 3. The result of this mapping is a monadic computation involving the  
91 Int, Arith, Read, and Let effects.

$$\begin{aligned}
\llbracket int(n) \rrbracket &= int\ n \\
\llbracket add(e_1, e_2) \rrbracket &= \llbracket e_1 \rrbracket \gg= \lambda x. \llbracket e_2 \rrbracket \gg= \lambda y. add\ x\ y \\
\llbracket sub(e_1, e_2) \rrbracket &= \llbracket e_1 \rrbracket \gg= \lambda x. \llbracket e_2 \rrbracket \gg= \lambda y. sub\ x\ y \\
\llbracket neg(e) \rrbracket &= \llbracket e \rrbracket \gg= \lambda x. neg\ x \\
\llbracket read \rrbracket &= read \\
\llbracket let(e, f) \rrbracket &= let\ \llbracket e \rrbracket\ (\lambda x. \llbracket f\ x \rrbracket) \\
\llbracket var(x) \rrbracket &= return\ x
\end{aligned}$$

■ **Figure 3** A denotational mapping from our abstract syntax onto our initial set of effectful operations.

92 Now that we have denoted our syntactic elements into our semantic domain as operations,  
93 we can start refining these operations to get closer to the desired target language which is  
94 X86 in our case. In the practice of algebraic effects, this refinement is facilitated by handlers  
95 as introduced by Plotkin and Pretnar [6]. These handlers give us access to the operations  
96 that occur in the program and the continuation of the program which we will simply call  $k$ .  
97 We only have to provide the rules that map the operation and continuation onto our semantic  
98 domain consisting of primitives, existing operations, or newly introduced operations.

## 23:4 Towards Modular Compilation using Higher-order Effects

The effect that we choose to handle first is the Let effect, which only has the *let* operation. We handle this operation by running the right hand side of the binding, passing the resulting value to the body, and finally passing the result of that to the continuation. In code that looks as follows:

```
handle (let e f) k = e >>= λx. f x >>= λz. k z
```

Note that this defines a strict semantics for our let bindings. By using a different handler we could give different semantics to our language. This is an example of the flexibility of algebraic effects and handlers.

effect X86 where	effect X86Var where
<i>addq</i> : <i>val</i> → <i>val</i> → <i>m</i> ()	<i>x86var</i> : <i>m val</i>
<i>subq</i> : <i>val</i> → <i>val</i> → <i>m</i> ()	
<i>negq</i> : <i>val</i> → <i>m</i> ()	
<i>movq</i> : <i>val</i> → <i>val</i> → <i>m</i> ()	
<i>callq</i> : <i>lab</i> → <i>m</i> ()	
<i>reg</i> : Register → <i>m val</i>	
<i>deref</i> : Register → $\mathbb{Z}$ → <i>m val</i>	

■ **Figure 4** The effects related to X86 and their operations with type signatures.

At this point, since our language is so simple we can already begin translating into our target language. In Figure 4 we show the a minimal subset of X86 that we need to compile our Arith and Read effects. This subset contains in-place arithmetic instructions, the ubiquitous move instruction, the call instruction, and an operation to inject concrete registers into our abstract value type. Additionally, we add an operation to generate fresh variables and inject them into our abstract value type.

We can translate our Arith effect operations into X86 operations by creating a fresh X86 variable, populating it, and then applying the in-place arithmetic operation to the variable. So, we write handler as follows:

```
handle (add x y) k = x86var >>= λz. movq x z >> addq y z >> k z
handle (sub x y) k = x86var >>= λz. movq x z >> subq y z >> k z
handle (neg x) k = x86var >>= λz. movq x z >> negq z >> k z
```

The *read* operation requires us to call a function that we will assume is defined in a standard library called `_read_int`. This function places its output in the `%rax` register, so we have to move it to avoid it being overwritten by other parts of our program. The full definition of our handler for the Read effect is as follows:

```
handle read k = x86var >>= λz. callq _read_int >> reg %rax >>= λx. movq x z >> k z
```

The final challenge to complete this minimal compiler pipeline is to allocate the X86 variables on the stack. Conceptually, this requires us to give each *x86var* operation and give each its own location of the stack. Keeping track of such information in our handler, however, is something we have not yet needed to do for the passes up to this point. Until now, we have handled each effect by translating into other effects directly. Instead, we can parameterize our handlers which means we pass along an extra parameter while handling

our operations. Parameterized handlers take one extra parameter and need to pass one extra argument to the continuation<sup>2</sup>. Now we can write the parameterized handler for the X86Var effect which assigns each variable to its own stack location as follows:

```
136   handle x86var k n = deref %rbp (-8 · n) >>= λz. k z (n + 1)
```

Note that we assume sufficient space is allocated on the stack. Additionally, when applying this handler we need to provide the starting value of the parameter *n*, which we will choose to be 1.

At this point, we have a full compiler pipeline from our source language to a subset of X86, but is still in the form of an effectful computation. To get a concrete representation, we implement two handlers for the remaining Int and X86 effects to produce an output string. As part of choosing this concrete representation, we also choose the concrete type for the variables *val* and *lab* to be the string type. We define the handler that turns our effectful computation of Int and X86 effects into a concrete string representation as follows:

```
147   handle (int n)      k = k (showInt n)
148   handle (addq x y)   k = "addq " ++ x ++ ", " ++ y ++ "\n" ++ k ()
149   handle (subq x y)   k = "subq " ++ x ++ ", " ++ y ++ "\n" ++ k ()
150   handle (negq x)     k = "negq " ++ x ++ "\n" ++ k ()
151   handle (movq x y)   k = "movq " ++ x ++ ", " ++ y ++ "\n" ++ k ()
152   handle (callq l)    k = "callq " ++ l ++ "\n" ++ k ()
153   handle (reg r)      k = k (showReg r)
154   handle (int r n)    k = k (showInt n ++ "(" ++ showReg r ++ ")")
```

### 3 Related Work

The origins of this work can be traced back to Eelco Visser's work on the Spoofox Language Workbench [9]. As part of Spoofox, Eelco is one of the designers of the Stratego [8] program transformation language. While Stratego can be used for developing compilers, Eelco was still looking for a way of specifying compilers that also abstracts over the control flow. We hope that this work can be the start of an answer to that research direction.

As mentioned in the introduction, our approach embraces the nanopass architecture [7, 4]. We improve upon this work by putting it on more formal foundations, making it fully typed to prevent common errors and even check that all cases are covered, and abstracting over the control flow in the compiler.

Our semantics-driven approach using an effect system is inspired by the work on symbolic execution by Wei et al. [10].

For our vision on verification of compilers we have taken inspiration from Interaction Trees [11]. Interaction Trees use algebraic laws of effects to prove the correctness of compilers. We hope to learn from that technique to prove the correctness of compilers written using our approach.

<sup>2</sup> We ignore effectful subcomputations, because they were already removed in an earlier pass.

## 172      4      Conclusions and Future Work

173      We have presented a new semantics-driven approach to writing compilers by using effect  
 174      operations as an intermediate representation. We use effect handlers to iteratively refine  
 175      operations in terms of increasingly lower level operations to finally reach a target machine  
 176      language.

177      We have shown a concrete example of this approach applied to a very simple language  
 178      with arithmetic and let-bound variables. We show the implementation of a denotation  
 179      function and handlers which compile this language is compiled in several passes to X86  
 180      machine language with variables. Currently, we are working on a stack allocation pass to  
 181      complete this minimal compiler.

182      In the future, we would like to complete the minimal compiler and extend it with more  
 183      complicated language constructs such as conditionals and anonymous functions. Additionally,  
 184      we would like to implement more complicated analyses on this effectful representation,  
 185      such as register allocation. We expect these analyses to consist of two stages: first derive  
 186      concrete structures such as control-flow graphs and interference graphs from our effectful  
 187      representation, and then perform a pass over the effectful computation that uses the results  
 188      of the analysis over these structures to transform the program. The first stage would be  
 189      similar to our handler that turns the effectful computation into a concrete string and the  
 190      second stage would employ a parameterized handler similar to our stack allocation handler.

191      Furthermore, we would like to explore the verification of our compilers using algebraic  
 192      laws for our effect operations. To be specific, we can define a set of laws that describe  
 193      the behavior of each of the effects in our compiler pipeline. If these laws are sound and  
 194      complete, with respect to for example definitional interpreters for the effects, then we can  
 195      prove compiler correctness by proving that these laws are preserved by each of our handlers.

## 196      —      References      —

- 197      1      Casper Bach Poulsen and Cas van der Rest. Hefty algebras – modular elaboration for  
 198      higher-order algebraic operations. Under submission, 2022.
- 199      2      Ana Bove, Peter Dybjer, and Ulf Norell. A brief overview of agda – a functional language  
 200      with dependent types. In Stefan Berghofer, Tobias Nipkow, Christian Urban, and Makarius  
 201      Wenzel, editors, *Theorem Proving in Higher Order Logics*, pages 73–78, Berlin, Heidelberg,  
 202      2009. Springer Berlin Heidelberg.
- 203      3      Adam Chlipala. Parametric higher-order abstract syntax for mechanized semantics. *SIGPLAN*  
 204      *Not.*, 43(9):143–156, sep 2008. URL: [https://doi-org.tudelft.idm.oclc.org/10.1145/](https://doi-org.tudelft.idm.oclc.org/10.1145/1411203.1411226)  
 205      1411203.1411226, doi:10.1145/1411203.1411226.
- 206      4      Andrew W. Keep and R. Kent Dybvig. A nanopass framework for commercial compiler  
 207      development. *SIGPLAN Not.*, 48(9):343–350, sep 2013. URL: [https://doi-org.tudelft.](https://doi-org.tudelft.idm.oclc.org/10.1145/2544174.2500618)  
 208      [idm.oclc.org/10.1145/2544174.2500618](https://doi-org.tudelft.idm.oclc.org/10.1145/2544174.2500618), doi:10.1145/2544174.2500618.
- 209      5      Gordon Plotkin and John Power. Adequacy for algebraic effects. In Furio Honsell and Marino  
 210      Miculan, editors, *Foundations of Software Science and Computation Structures*, pages 1–24,  
 211      Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- 212      6      Gordon Plotkin and Matija Pretnar. Handlers of algebraic effects. In Giuseppe Castagna,  
 213      editor, *Programming Languages and Systems*, pages 80–94, Berlin, Heidelberg, 2009. Springer  
 214      Berlin Heidelberg.
- 215      7      Dipanwita Sarkar, Oscar Waddell, and R. Kent Dybvig. A nanopass infrastructure for  
 216      compiler education. In *Proceedings of the Ninth ACM SIGPLAN International Conference on*  
 217      *Functional Programming*, ICFP '04, page 201–212, New York, NY, USA, 2004. Association for  
 218      Computing Machinery. URL: [https://doi-org.tudelft.idm.oclc.org/10.1145/1016850.](https://doi-org.tudelft.idm.oclc.org/10.1145/1016850.1016878)  
 219      1016878, doi:10.1145/1016850.1016878.

- 220   8   Eelco Visser. Stratego: A language for program transformation based on rewriting strategies  
221       system description of stratego 0.5. In Aart Middeldorp, editor, *Rewriting Techniques and*  
222       *Applications*, pages 357–361, Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- 223   9   Guido H. Wachsmuth, Gabriël D.P. Konat, and Eelco Visser. Language design with the  
224       spoofax language workbench. *IEEE Software*, 31(5):35–43, 2014. doi:10.1109/MS.2014.100.
- 225  10   Guannan Wei, Oliver Bračevac, Shangyin Tan, and Tiark Rumpf. Compiling symbolic execution  
226       with staging and algebraic effects. *Proc. ACM Program. Lang.*, 4(OOPSLA), nov 2020. URL:  
227       <https://doi-org.tudelft.idm.oclc.org/10.1145/3428232>, doi:10.1145/3428232.
- 228  11   Li-yao Xia, Yannick Zakowski, Paul He, Chung-Kil Hur, Gregory Malecha, Benjamin C. Pierce,  
229       and Steve Zdancewic. Interaction trees: Representing recursive and impure programs in coq.  
230       *Proc. ACM Program. Lang.*, 4(POPL), dec 2019. URL: <https://doi-org.tudelft.idm.oclc.org/10.1145/3371119>, doi:10.1145/3371119.  
231