

# Towards Modular Compilation using Higher-order Effects

Anonymous author

Anonymous affiliation

## Abstract

Compilers transform a human readable source language into machine readable target language. Nanopass compilers simplify this approach by breaking up this transformation into small steps that are more understandable, maintainable, and extensible. We propose a semantics-driven variant of the nanopass compiler architecture exploring the use of effects and handlers to model the intermediate languages and the transformation passes, respectively. Our approach is fully typed and ensures that all cases in the compiler are covered. Additionally, by using an effect system we abstract over the control flow of the intermediate language making the compiler even more flexible. We apply this approach to a minimal compiler from a language with arithmetic and let-bound variables to a minimal subset of X86 machine code with variables. In the future, we hope to extend this work to compile a larger and more complicated language and we envision a formal verification framework from compilers written in this style.

**2012 ACM Subject Classification** Theory of computation → Program semantics; Software and its engineering → Compilers

**Keywords and phrases** algebraic effects and handlers, higher-order effects, monadic semantics, modularity, compilation, nanopass

**Digital Object Identifier** 10.4230/OASIS.CVIT.2016.23

## 1 Introduction

The essence of a compiler is a function from a source language, suitable for humans, to a machine language, suitable for computers. As our computers have become more powerful we have seen increasingly complex compilers providing extensive safety guarantees and powerful optimizations. To manage this complexity, modern compilers are designed as a composition of multiple passes. However, the total number of passes has traditionally been kept low for the sake of performance, because each pass adds extra overhead. Thus, compiler passes are more complicated than necessary and therefore harder to *understand*, *maintain*, and *extend*.

To address this problem, Sarkar, Waddell, and Dybvig [7] introduce the nanopass compiler architecture. In the nanopass architecture, each pass is designed to be as simple as possible. It is not a problem to use many more passes than are used in traditional compilers. To address concerns about the performance of this architecture, Keep and Dybvig [4] show that it is possible to write a competitive commercial compiler using this architecture.

While the development of the nanopass architecture and the development of a commercial compiler is a great engineering achievement, we believe the theoretical foundation is underexplored. In this paper, we present our ongoing work on developing a semantics-driven nanopass compiler architecture. Our approach has the practical advantage of preventing type errors in the compiler and ensuring all cases are covered. Additionally, our approach abstracts over the control flow giving many of the benefits of continuation-passing-style while retaining a simple monadic interface. During the development of our approach, we are anticipating the possibility of verifying the correctness of compilers written in this style.

Concretely, we make the following contributions:

- We introduce a novel approach to designing and implementing practical compilers while staying close to formal denotational semantics specifications (Section 2).

■ We demonstrate our approach on a simple language with arithmetic and let-bound variables (Section 2). We compile this language to a subset of X86 with variables. We are still working a method to allocate these variables on the stack.

## 2      Compiling with Higher-order Effects

In this section, we present our approach by applying it to a very simple language with arithmetic and let-bound variables. The target language of our compiler is X86 machine code. We explain the required concepts as we develop our compiler for this language. The specifications and compiler passes are presented in a simplified notation, but we have implemented all the work we present here in the Agda programming language [2].

We start off by assuming our parser and possibly type checker has finished and produced an abstract syntax tree which follows the grammar described in Figure 1. Our language has integers, addition, subtraction, negation, a read operation to read an input integer, and a let to bind variables and a var to refer to bound variables. The abstract syntax is unusual because we reuse the variable binding facilities from our host language in the form of the  $\lambda x$ . binding in the let constructor. This style of abstract syntax is called parametric higher-order abstract syntax (PHOAS) [3]. It allows us to avoid the complexities of variable binding and thus simplify our presentation.

$$\begin{aligned} \text{expr} ::= & \text{int}(n) \\ & | \text{add}(\text{expr}, \text{expr}) \\ & | \text{sub}(\text{expr}, \text{expr}) \\ & | \text{neg}(\text{expr}) \\ & | \text{read} \\ & | \text{let}(\text{expr}, \lambda x. \text{expr}) \\ & | \text{var}(x) \end{aligned}$$

■ **Figure 1** Abstract syntax of our simple language with arithmetic and let-bound variables.

The first step of our compilation pipeline will be to denote these syntactic constructs onto a set of semantic algebraic operations in the sense of algebraic effects [5]. As is customary when using algebraic effects in functional programming languages, we group these operations under units we call effects. We could group every operation under a single effect, however with the benefit of hindsight we decide to distribute the operations over four effects: Int, Arith, Read, and Let.

Figure 2 shows the operations that correspond directly to our source language. However, there are some particularities we address individually:

- To keep our example simple, we have chosen to use this single type for all values, but we keep the type abstract and simply call it ‘val’. It is crucial to keep this type abstract for reasons we explain at the end of this section. Also note that we now need a special *int* operation to inject integers into this abstract value type.
- We write the surrounding monadic context as  $m$ , which provides the standard  $\gg=$ ,  $\gg$ , and *return* operations. The monadic context of each operation always includes at least the effect that the operation belongs to, but it can accomodate other effects too. For

$$\begin{array}{ll}
\text{int} : \mathbb{Z} \rightarrow m \text{ val} & \\
& \text{read} : m \text{ val} \\
\\
\text{add} : \text{val} \rightarrow \text{val} \rightarrow m \text{ val} & \\
\text{sub} : \text{val} \rightarrow \text{val} \rightarrow m \text{ val} & \text{let} : m \text{ val} \rightarrow (\text{val} \rightarrow m \text{ val}) \rightarrow m \text{ val} \\
\text{neg} : \text{val} \rightarrow m \text{ val} &
\end{array}$$

■ **Figure 2** The operations of our source language with their signatures. Divided into the following effects from left to right and then top to bottom: Int, Read, Arith, and Let.

78 example, the monadic context of  $\text{int } 1 \gg\gg \lambda x. \text{add } x \ x$  contains at least the Int and Arith  
 79 effects.

80 ■ Readers knowledgeable about effect systems might notice that the *let* operation has  
 81 arguments that are themselves monadic computations. In standard algebraic effects  
 82 and handlers this is not allowed, however our approach uses a novel higher-order effect  
 83 formalism that does support such effectful subcomputations [1].

84 The first pass of our compiler pipeline is to map our abstract syntax from Figure 1 onto  
 85 the operations we have defined for our source language from Figure 2. This mapping, called  
 86 a denotation and written using the  $\llbracket \cdot \rrbracket$  notation, is a recursive traversal of the abstract syntax  
 87 tree shown in Figure 3. The result of this mapping is a monadic computation involving the  
 88 Int, Arith, Read, and Let effects.

$$\begin{aligned}
\llbracket \text{int}(n) \rrbracket &= \text{int } n \\
\llbracket \text{add}(e_1, e_2) \rrbracket &= \llbracket e_1 \rrbracket \gg\gg \lambda x. \llbracket e_2 \rrbracket \gg\gg \lambda y. \text{add } x \ y \\
\llbracket \text{sub}(e_1, e_2) \rrbracket &= \llbracket e_1 \rrbracket \gg\gg \lambda x. \llbracket e_2 \rrbracket \gg\gg \lambda y. \text{sub } x \ y \\
\llbracket \text{neg}(e) \rrbracket &= \llbracket e \rrbracket \gg\gg \lambda x. \text{neg } x \\
\llbracket \text{read} \rrbracket &= \text{read} \\
\llbracket \text{let}(e, f) \rrbracket &= \text{let } \llbracket e \rrbracket (\lambda x. \llbracket f \rrbracket x) \\
\llbracket \text{var}(x) \rrbracket &= \text{return } x
\end{aligned}$$

■ **Figure 3** Denotational mapping from our abstract syntax onto our initial set of effectful operations.

89 Now that we have denoted our syntactic elements into our semantic domain as operations,  
 90 we can start refining these operations to get closer to the desired target language which is  
 91 X86 in our case. In the practice of algebraic effects, this refinement is facilitated by handlers  
 92 as introduced by Plotkin and Pretnar [6]. These handlers give us access to the operations  
 93 that occur in the program and the continuation of the program which we will simply call  $k$ .  
 94 We only have to provide the rules that map the operation and continuation onto our semantic  
 95 domain consisting of primitives, existing operations, or newly introduced operations.

96 The effect that we choose to handle first is the Let effect, which only has the *let* operation.  
 97 We handle this operation by running the right hand side of the binding, passing the resulting  
 98 value to the body, and finally passing the result of that to the continuation. In code that

## 23:4 Towards Modular Compilation using Higher-order Effects

99 looks as follows:

100     handle (let  $e\ f$ )  $k = e \gg \lambda x. f\ x \gg \lambda z. k\ z$   
 101

102 Note that this defines a strict semantics for our let bindings. By using a different handler  
 103 we could give different semantics to our language. This is an example of the flexibility of  
 104 algebraic effects and handlers.

$addq : val \rightarrow val \rightarrow m\ ()$	
$subq : val \rightarrow val \rightarrow m\ ()$	
$negq : val \rightarrow m\ ()$	
$movq : val \rightarrow val \rightarrow m\ ()$	$x86var : m\ val$
$callq : lab \rightarrow m\ ()$	
$reg : Register \rightarrow m\ val$	

■ **Figure 4** X86 related effects. On the left X86. On the right X86Var.

105 At this point, since our language is so simple we can already begin translating into  
 106 our target language. In Figure 4 we show the a minimal subset of X86 that we need to  
 107 compile our Arith and Read effects. This subset contains in-place arithmetic instructions,  
 108 the ubiquitous move instruction, the call instruction, and an operation to inject concrete  
 109 registers into our abstract value type. Additionally, we add an operation to generate fresh  
 110 variables and inject them into our abstract value type.

111 We can translate our Arith effect operations into X86 operations by creating a fresh X86  
 112 variable, populating it, and then applying the in-place arithmetic operation to the variable.  
 113 So, we write handler as follows:

114     handle (add  $x\ y$ )  $k = x86var \gg \lambda z. movq\ x\ z \gg addq\ y\ z \gg k\ z$   
 115     handle (sub  $x\ y$ )  $k = x86var \gg \lambda z. movq\ x\ z \gg subq\ y\ z \gg k\ z$   
 116     handle (neg  $x$ )  $k = x86var \gg \lambda z. movq\ x\ z \gg negq\ z \gg k\ z$   
 117

118 The *read* operation requires us to call a function that we will assume is defined in a standard  
 119 library called `read_int`. This function places its output in the `%rax` register, so we have to  
 120 move it to avoid it being overwritten by other parts of our program. The full definition of  
 121 our handler for the Read effect is as follows:

122     handle read  $k = x86var \gg \lambda z. callq\ read\_int \gg reg\ \%rax \gg \lambda x. movq\ x\ z \gg k\ z$   
 123

124 The final challenge to complete this minimal compiler is to allocate the X86 variables  
 125 on the stack. Conceptually, this requires us to give each *x86var* operation and give each its  
 126 own location of the stack. Keeping track of such information in our handler, however, is  
 127 something we have not yet needed to do for the passes up to this point. The main problem is  
 128 that we need to pass a valid value as input to the continuation of each operation to get access  
 129 to the operations in the continuation. When the type of the argument of the continuation is a  
 130 unit type, such as for the X86 arithmetic, move, and call instructions, then we can simply  
 131 construct the unit value and pass that to the continuation. However, if we allow arbitrary  
 132 concrete types as results of our operations then there is no guarantee that we can construct a  
 133 value of that type and it is not always possible to reconstruct the higher-order representation  
 134 if we do manage to construct such a value.

135        Luckily, we have chosen to keep our value type abstract. So, we still have a choice to  
136        instantiate it to a type that suits our purpose. Instantiating it to the unit type would make  
137        it possible for us to construct a value to pass to the continuation. However, we would no  
138        longer be able to distinguish values passed to different continuations, so it would not be  
139        possible to reconstruct a higher order representation for further manipulation. Instead, our  
140        solution is to instantiate it to the type of natural numbers and to pass a unique number to  
141        each continuation. Whenever we encounter such a natural number in the rest of the program,  
142        we know which continuation it originated from. Hence, we are able to reconstruct a higher  
143        order operation.

144        This process sounds complicated, but we expect it is possible to expose the ability to  
145        keep track of the required information through an easy to use API.

### 146    **3    Related Work**

147        The origins of this work can be traced back to Eelco Visser's work on the Spoofox Language  
148        Workbench [9]. As part of Spoofox, Eelco is one of the designers of the Stratego [8] program  
149        transformation language. While Stratego can be used for developing compilers, Eelco was  
150        still looking for a way of specifying compilers that also abstracts over the control flow. We  
151        hope that this work can be the start of an answer to that research direction.

152        As mentioned in the introduction, our approach embraces the nanopass architecture [7, 4].  
153        We improve upon this work by putting it on more formal foundations, making it fully typed  
154        to prevent common errors and even check that all cases are covered, and abstracting over  
155        the control flow in the compiler.

156        Our semantics-driven approach using an effect system is inspired by the work on symbolic  
157        execution by Wei et al. [10].

158        For our vision on verification of compilers we have taken inspiration from Interaction  
159        Trees [11]. Interaction Trees use algebraic laws of effects to prove the correctness of compilers.  
160        We hope to learn from that technique to prove the correctness of compilers written using our  
161        approach.

### 162    **4    Conclusions and Future Work**

163        We have presented a new semantics-driven approach to writing compilers by using effect  
164        operations as an intermediate representation. We use effect handlers to iteratively refine  
165        operations in terms of increasingly lower level operations to finally reach a target machine  
166        language.

167        We have shown a concrete example of this approach applied to a very simple language  
168        with arithmetic and let-bound variables. We show the implementation of a denotation  
169        function and handlers which compile this language is compiled in several passes to X86  
170        machine language with variables. Currently, we are working on a stack allocation pass to  
171        complete this minimal compiler.

172        In the future, we would like to complete the minimal compiler and extend it with more  
173        complicated language constructs such as conditionals and anonymous functions. Furthermore,  
174        we would like to explore the verification of our compilers using algebraic laws for our effect  
175        operations.

---

References

---

- 177    **1**    Casper Bach Poulsen and Cas van der Rest. Hefty algebras – modular elaboration for  
178       higher-order algebraic operations. Under submission, 2022.
- 179    **2**    Ana Bove, Peter Dybjer, and Ulf Norell. A brief overview of agda – a functional language  
180       with dependent types. In Stefan Berghofer, Tobias Nipkow, Christian Urban, and Makarius  
181       Wenzel, editors, *Theorem Proving in Higher Order Logics*, pages 73–78, Berlin, Heidelberg,  
182       2009. Springer Berlin Heidelberg.
- 183    **3**    Adam Chlipala. Parametric higher-order abstract syntax for mechanized semantics. *SIGPLAN*  
184       *Not.*, 43(9):143–156, sep 2008. URL: [https://doi-org.tudelft.idm.oclc.org/10.1145/](https://doi-org.tudelft.idm.oclc.org/10.1145/1411203.1411226)  
185       [1411203.1411226](https://doi-org.tudelft.idm.oclc.org/10.1145/1411203.1411226), doi:10.1145/1411203.1411226.
- 186    **4**    Andrew W. Keep and R. Kent Dybvig. A nanopass framework for commercial compiler  
187       development. *SIGPLAN Not.*, 48(9):343–350, sep 2013. URL: [https://doi-org.tudelft.](https://doi-org.tudelft.idm.oclc.org/10.1145/2544174.2500618)  
188       [idm.oclc.org/10.1145/2544174.2500618](https://doi-org.tudelft.idm.oclc.org/10.1145/2544174.2500618), doi:10.1145/2544174.2500618.
- 189    **5**    Gordon Plotkin and John Power. Adequacy for algebraic effects. In Furio Honsell and Marino  
190       Miculan, editors, *Foundations of Software Science and Computation Structures*, pages 1–24,  
191       Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- 192    **6**    Gordon Plotkin and Matija Pretnar. Handlers of algebraic effects. In Giuseppe Castagna,  
193       editor, *Programming Languages and Systems*, pages 80–94, Berlin, Heidelberg, 2009. Springer  
194       Berlin Heidelberg.
- 195    **7**    Dipanwita Sarkar, Oscar Waddell, and R. Kent Dybvig. A nanopass infrastructure for  
196       compiler education. In *Proceedings of the Ninth ACM SIGPLAN International Conference on*  
197       *Functional Programming*, ICFP ’04, page 201–212, New York, NY, USA, 2004. Association for  
198       Computing Machinery. URL: [https://doi-org.tudelft.idm.oclc.org/10.1145/1016850.](https://doi-org.tudelft.idm.oclc.org/10.1145/1016850.1016878)  
199       [1016878](https://doi-org.tudelft.idm.oclc.org/10.1145/1016850.1016878), doi:10.1145/1016850.1016878.
- 200    **8**    Eelco Visser. Stratego: A language for program transformation based on rewriting strategies  
201       system description of stratego 0.5. In Aart Middeldorp, editor, *Rewriting Techniques and*  
202       *Applications*, pages 357–361, Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- 203    **9**    Guido H. Wachsmuth, Gabriël D.P. Konat, and Eelco Visser. Language design with the  
204       spoofax language workbench. *IEEE Software*, 31(5):35–43, 2014. doi:10.1109/MS.2014.100.
- 205    **10**    Guannan Wei, Oliver Bračevac, Shangyin Tan, and Tiark Rompf. Compiling symbolic execution  
206       with staging and algebraic effects. *Proc. ACM Program. Lang.*, 4(OOPSLA), nov 2020. URL:  
207       <https://doi-org.tudelft.idm.oclc.org/10.1145/3428232>, doi:10.1145/3428232.
- 208    **11**    Li-yao Xia, Yannick Zakowski, Paul He, Chung-Kil Hur, Gregory Malecha, Benjamin C. Pierce,  
209       and Steve Zdancewic. Interaction trees: Representing recursive and impure programs in coq.  
210       *Proc. ACM Program. Lang.*, 4(POPL), dec 2019. URL: [https://doi-org.tudelft.idm.oclc.](https://doi-org.tudelft.idm.oclc.org/10.1145/3371119)  
211       [org/10.1145/3371119](https://doi-org.tudelft.idm.oclc.org/10.1145/3371119), doi:10.1145/3371119.